

Multi-pose people detection in 3D point clouds

Stefano Zanella - 621796

Feb 2017

1 Abstract

This paper describes a people detection algorithm for three-dimensional point clouds based on the well-known Viola-Jones face recognition algorithm [2]. The key contribution of the project is an implementation of mentioned algorithm, adapted for work with 3D point clouds. The implementation consists in a cascaded classifier and its training algorithm. This detector serves as a cornerstone for the overall people detection algorithm, which features a variant of the region-growing segmentation algorithm found in the PCL library. Analysis of the algorithm's performance are provided based on a given set of 40 images.

2 Introduction

This paper attempts at extending the current research in the field of recognition of human beings in three-dimensional point clouds. The key intuition behind the proposed approach is that human presence in a scene might be connected to the presence of any part of a human body.

Building a detector flexible enough to recognize any body part (e.g. a face or a hand), or any sensible combination of them (e.g. a leg and a foot), is a task that has low chances to lead to good performances. Indeed, also designing a detector that recognizes bodies in multiple poses is not a simple task, as the flexibility of human body makes the number of possible poses very high.

Moreover, such a detector might not perform well in case of obstructed bodies or bodies only partially framed into the scene.

Therefore, the proposed approach starts from the most distinguishing part of a human body: its head. A human head, put in the context of a larger scene, provides some desirable features:

- it's relatively rigid, meaning that it acts approximately as any other rigid body that can only be subject to rotation. This is important as it allows to be detected by using standard object recognition techniques
- it's small compared to the rest of the body, opening the door to algorithms that achieve higher performances than those than need to perform larger scans of a scene or analyze larger sets of points at a single time
- it's possibly the body part that's most discriminative of human nature: e.g. it might be hard, at a relative distance, to say if a barefoot belongs to a dummy; on the other hand, it's definitely easier to discriminate if a frontal head it's human or not. This might be of help to recognition algorithms.

The subsequent reasoning step is that, given an accurate way of detecting a human head in a scene, we can extend the detection to points that seem to belong to the same object as the head. This is a task that segmentation algorithms can solve pretty well, with a good range of choices in terms of speed, accuracy and performances.

This "detection expansion" technique offers a few advantages:

- it allows to detect bodies that are partially obstructed or not entirely in the scene
- it makes the algorithm robust to changes of pose, given the segmentation algorithm can correctly identify "compact" objects (i.e. clusters

of points that are mostly equidistant from each other).

- if decomposes a "hard" problem into two problems that have been already well treated in literature, and for which a number of efficient solutions exist.

As a matter of implementation, two well-known algorithms have been selected:

- for head (face) detection, the algorithm devised by Viola and Jones [2] has been selected and implemented with a number of adaptations and improvements made possible by its application in the 3D space
- for body detection, a variant of the region-growing segmentation algorithm present in the PCL library has been used

Overview

The remaining of this paper will describe more in detail the people detection algorithm implementation and its peculiarities. We will start by describing the face detection algorithm and all the specific changes that have been introduced for it to be applicable to the 3D space. After that the text will focus on the segmentation algorithm. Finally the results of running the algorithm described on a validation set of 40 images will be discussed. A selection of ideas for future improvements closes the paper.

3 Face detection

The first step in the processing chain is to detect faces in the current scene. This is done by employing a boosted cascade detector that takes heavy inspiration from the work presented by Viola and Jones in their seminal paper. The detector is composed of a set of rectangular features that get evaluated on the area currently analyzed, then combining the results in a linear combination and evaluating the final result against a threshold.

Each linear combination of features represents a stage in the final detector: if an area is recognized as a face in a given stage, the algorithm moves to the next stage. The algorithm terminates when a stage rejects a sample, or when all the stages accept it.

3.1 Features

A feature is a function over a rectangular area of the sample being analyzed. In general terms, a feature defines a linear combination that reduces the

value of the pixels in the sample to a single number. More specifically, features can be of 3 types:

- two-rectangle features: they are defined as the difference between the sum of the pixels in two rectangular areas. The areas have same size, same shape and are horizontally or vertically adjacent
- three-rectangle features: in this case what gets computed is the sum of two outside rectangles, subtracted from the sum of an inner rectangle of same shape and size as the other two
- four-rectangle features: the rectangles are positioned in a grid and the features computes the difference between the sum of diagonal pairs

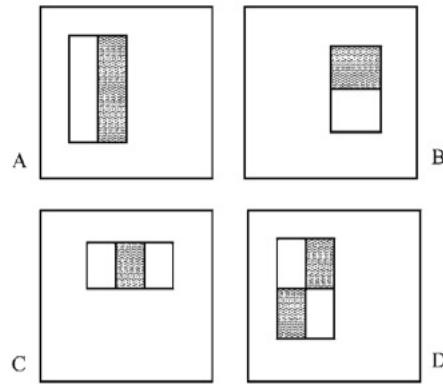


Figure 1: Examples of features relative to the detection window

In order to make the implementation presented in this paper as close as possible to the original algorithm described by Viola and Jones, features have been calculated based only on the 2D RGB information of the point cloud. This makes features calculation very fast and easy to reasonate about as the problem can be modeled with simple two-dimensional array access. This means the point cloud passed as input to the detector must be organized.

3.2 Cascaded classifier

Evaluation of a single feature is hardly sufficient to obtain a result that can be distinguished by a random classifier operating over a uniform probability distribution. A classifier operating with a single feature is named for this reason *weak classifier*.

Given a detection window x , a feature f , a threshold θ and a polarity p , a weak classifier can be described by the equation:

$$h(x, f, p, \theta) = \begin{cases} 1, & \text{if } pf(x) < p\theta \\ 0, & \text{otherwise} \end{cases}$$

To obtain more robust results, weak classifiers can be combined together in a linear combination. Such a combination is referred to as *strong classifier*. By means of a strong classifier it's possible to obtain results that are satisfactory in terms of detection rate and false positive rate, but at the expense of an increased need for computational power.

The general equation of a strong classifier composed of T weak classifiers is defined as follows:

$$C(X) = \begin{cases} 1 & \sum_{t=1}^T \log \frac{1 - \epsilon_t}{\epsilon_t} h_t(X) \geq \frac{1}{2} \sum_{t=1}^T \log \frac{1 - \epsilon_t}{\epsilon_t} \\ 0 & \text{otherwise} \end{cases}$$

where

- $h_n(X)$ is the n th weak classifier, defined as above
- ϵ_n is the weighted error for the n th classifier, calculated as part of the training process

The training process for a strong classifier starts from a set of training images (x_i, y_i) pre-classified as positive ($y_i = 1$) or negative ($y_i = 0$), and proceeds as follows:

- initialize weights $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$ for $y_i = 0, 1$ respectively, where m and l are the number of negative and positive samples
- while the false positive rate is above the target level:
 - normalize the weights

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

- select the best weak classifier with respect to the weighted error

$$\epsilon_t = \min_{f,p,\theta} \sum_i w_i |h(x_i, f, p, \theta) - y_i|$$

- define

$$h_t(x) = h(x, f_t, p_t, \theta_t)$$

where f_t, p_t, θ_t are the minimizers of ϵ_t

- update the weights

$$w_{t+1,i} = w_{t,i} \left(\frac{\epsilon_t}{1 - \epsilon_t} \right)^{1-e_i}$$

where $e_i = 0$ if example X_i is classified correctly, 1 otherwise

This approach would already allow to build a strong classifier that performs much better than a single weak classifier, but increasing the performance of the classifier directly impacts computation time.

This can be improved by leveraging a simple idea: in most cases, a detector is going to reject a sample (i.e. there aren't that many faces in a scene compared to the overall number of possible samples). So a simple classifier could already be useful in reducing the number of samples that need to be checked by a more complex (and thus refined) classifier.

This approach is achieved by implementing a variant of AdaBoost [1] that connects a series of strong classifiers of varying complexity in a chain. Ideally the detectors should have high detection rate and acceptable false positive rate. That given, the cascade works as follows: if a sample is rejected in a stage of the cascade, it won't be processed by later stages. A sample is declared positive only if all the stages classify it so.

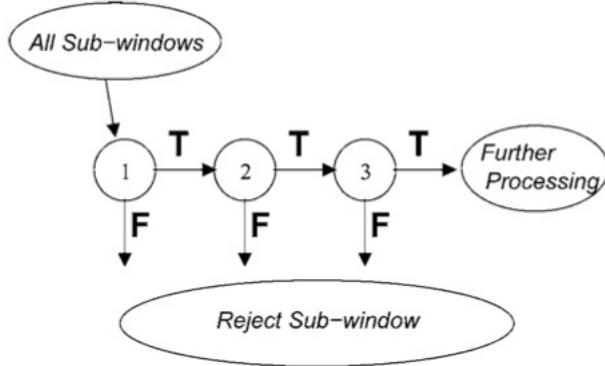


Figure 2: Schematic depiction of a cascade classifier

If we sort the classifiers in order of complexity we would greatly reduce the average computational power required to classify a sample. Only the actual faces will require to go through the whole cascade, but the vast majority of the samples will be rejected early.

The training algorithm for a cascade classifier builds on top of the one for training a strong classifier. It starts from a fixed target for the maximum acceptable false positive rate f , the overall target false positive rate F_{target} and for the minimum acceptable detection rate d , and proceeds as follows:

- $F_0 = 1.0, D_0 = 1.0$
- $i = 0$
- while $F_i > F_{target}$
 - $i++$
 - create a new stage in the classifier with a strong classifier
 - while $F_i > f \cdot F_{i-1}$
 - * train a new weak classifier to add to the last strong classifier
 - * evaluate current cascaded classifier on a validation set to determine F_i and D_i
 - * while $D_i < d$
 - decrease threshold of the last trained weak classifier
 - if $F_i > F_{target}$ evaluate current cascaded classifier on set of non-face images and replace the current negative set used for training with images that are wrongly detected

3.3 Integral image

For efficient computation of features (both in detection and training), an integral image representation of the original sample is computed.

In an integral image, each pixel is the sum of all the pixels above and to the left.

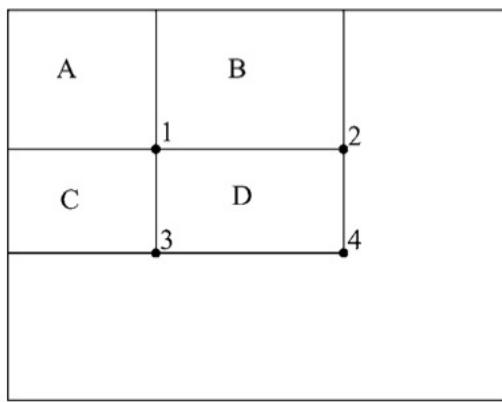


Figure 3: The sum of the pixels within rectangle D can be computed with four array references: $4 + 1 - 2 - 3$

Using this representation, the value of a rectangle area can be computed by 4 array references as illustrated in the picture above. Computing any of the features described in the previous paragraph thus requires a constant number of operations, independent from the size of the feature.

In practical terms, in order to perform all the required operations on an image efficiently, we need to store both the integral image and the integral image of the squares, that is the matrix in which each cell contains the square sum of the cells above and to the left. This is needed in order to calculate the image variance with $O(1)$ complexity.

Another interesting aspect of the implementation of the integral image is the different way it is accessed by the training algorithm and the detection algorithm. For the training algorithm all that's needed is to have an integral image of the entire sample. A specific feature is calculated only once on the overall sample. On the contrary, the detection algorithm needs to calculate features multiple times across the same image. More importantly, image normalization must happen on the overall sample in case of the training algorithm, whereas in the detector each of the image portions scanned needs to be normalized independently from the rest.

These two characteristics pose two problems: how to share the code that calculates the integral image between the two use cases, and how to make image normalization as efficient as possible.

The first problem has been solved by collecting the integral image and square integral image calculation code in a common class named `IntegralImage`. This class provides facilities to get various calculations for a specific rectangle of the image held by a class' instance, such as the area sum, mean and variance. Taking it from there, then, in order to solve the second problem two classes have been designed: `TrainingSample` and `SubWindow`. A training sample maps the use case scenario found in training, which boils down to simple area sum calculation using the whole integral image. A training sample also stores information useful for the training itself, as for example a flag saying if the sample contains a face or not (i.e. if it should be classified positively or not). A sub-window, instead, maps the use case found in the detector. The main feature is the possibility to calculate a feature on a portion of the image that gets normalized independently from the rest of the image.

Image normalization is required to make training and detection as independent as possible from lighting conditions and general camera exposure. A form of normalization useful in this case is variance normalization. This transforms the input so that the end result has a $N(0, 1)$ distribution. Given a pixel value x belonging to an image X , its variance-normalized counterpart can be expressed as:

$$x_{norm} = \frac{x - \bar{x}}{\sigma^2}$$

where \bar{x} is the image's mean value, defined as:

$$\bar{x} = \frac{1}{N} \sum X$$

whereas σ^2 is the square variance, defined as:

$$\sigma^2 = \bar{x}^2 - \frac{1}{2} \sum X^2$$

We can then calculate the linear combination of a set of variance-normalized pixels by relying on the following equation:

$$\sum X_{norm} = \frac{\sum X}{\sigma^2}$$

We need then to solve the problem of how to efficiently calculate the value of a feature normalized over a portion of the overall image. We do this by allowing a **SubWindow** to be repositioned inside the original image. When repositioning the sub-window, we recalculate the variance over the new image portion. This is a step that requires $O(1)$ since the integral square image is stored alongside with the integral image. At this point we can simply apply the last equation whenever we calculate an area sum.

To make things easier, an instance of **SubWindow** stores the coordinates of the image portion being analyzed (in the form of its top-left and bottom-right corner). When repositioning the sub-window, we also recalculate and store the variance of the image portion. By repositioning the sub-window instead of the feature itself, we can achieve a consistent and position-independent representation of a feature (i.e. the sub-window takes care of translating and scaling the feature's coordinates).

3.4 Sizing the features

The original implementation of the Viola-Jones algorithm is designed to work on 2D images. This means it doesn't accept any depth information. For this reason, in the original paper, the final detector needs to be run at different scales in each area in order to take into account the fact that faces might be present in different sizes.

In addition to that, a training algorithm has to focus on limited areas of a specific scene. We'll call these areas *samples*. The size of these samples is

also affected by the distance, and bears a direct linear relationship to the size of the features. Determination of the size of a sample obeys the same rules as the determination of a feature's size.

This extra computational step can be avoided if we consider depth information present in a point cloud. In particular we can start from a reference sample size, and then scale it based on the position on the Z-axis of the area of the scene currently analyzed.

The reference sample should be big enough to guarantee it fits heads of any size, but not any bigger. In other words, the reference sample should be as small as possible without running into the risk of "cutting off" heads that appear bigger because of camera distortion.

If the reference sample size is expressed for a sample lying on average 1m away from the camera, then the size of any given sample can then be calculated just by dividing the reference size by the sample distance.

In practical terms, this poses two problems: the fact that the calculation happens on pixels and not on real values, and that the depth is not constant over a specific area.

The first problem is simply solved by rounding up the division to the next integer. Moreover, since once the sample size has been calculated we need to calculate the precise pixel coordinates of the rectangle to consider, we take the extra step of rounding up the upper coordinates (i.e. bottom-right corner), while at the same time we round down the lower coordinates (i.e. top-left corner). This ensures that even at a big distance we minimize the effect of rounding by oversizing the sample square.

The problem of the distance not being constant over an area is solved by always considering the point being analyzed as the center of a rectangle. While this seems overly simplistic, we also leverage the sequential scanning of the image. Given that the reference sample is sized so that small variations in the final sample size don't affect the detector's performance, if at a certain point we encounter an outlier (either a NaN or a point with a lot of noise in the distance information), we can just assume that at least one of its neighbors contains correct information. In other words, in case we encounter outliers, we can safely assume that the algorithm already classified or will classify the area correctly when analyzing neighbor points.

The only downside of this approach is that in case of areas that are overall noisy (such as peripheral areas in the camera's field of view) the algorithm might end up skipping classification on large portions of the image. But in that case, classification is generally difficult anyway given the low SNR.

Another small inefficiency present in the scanning loop due to this approximation of an area's average distance is the computation that happens at the boundaries of the scene. Since a specific point is considered the center

of a square when determining the sample size and position, there must be enough points in all directions so that a full square can be extracted and analyzed. If the scanning algorithm is working near the image's boundaries, this condition might not be respected. In this case, the algorithm also skips detection, based on the assumption that at some point the same algorithm will analyze a close neighbor whose distance will allow to include all the points discarded up to that point because the sample would overflow.

The calculation of the size of a feature is performed on a similar basis but with a caveat: each feature needs to record additional information relative to the size of the sample that was used when training it. This is necessary because by introducing information about image depth in the training algorithm, a specific feature's size loses all meanings when the feature needs to be rescaled, if the same information is not related to the size of the sample it was trained with.

In addition to the point above, depth information is also used to minimize the amount of features generated and evaluated when training the classifier. In the 2D version of the algorithm, all features of all size from 1x1 to the size of the image must be generated. In the case of a 3D point cloud, we can reduce the amount of features to those that fit the smallest sample present in the scene. Features for bigger samples can be upscaled from the smaller ones. This approach offers two contrasting aspects:

- on one side, resolution is lost when upscaling the features (the bigger the size differences among samples, the bigger the loss)
- on the other side, starting from a set of features for a bigger sample, opens the door to the risk of downscaling a feature to the point it goes to a sub-pixel level, resulting in a complete loss of information during the training process.

Trials on the designed algorithm showed that the problem introduced by downsampling features to sub-pixel sizes was jeopardizing the algorithm's effectiveness, so the first option was selected.

4 Body segmentation

Once a face has been identified by the detection algorithm, an attempt is made to expand the detection and identify the points that belong to the cloud. The idea is that since we already know some of the points that belong to the positive classification, we can use a segmentation technique that starts from those points and tries to include all the points that supposedly belong to the same object. One such technique is already implemented in PCL and goes under the name of region growing.

Region growing is particularly suitable due to the fact that it requires initial seed points, and the classification algorithm is designed to provide those.

The main idea of the region growing algorithm is to start from a seed point and put all the points that match a certain definition of neighbor into the same region set. The curvature of each point is then checked against a certain predefined threshold: if it's below the limit, then the point is also added to the set of seeds.

Once the set of seeds have been exhausted, the algorithm considers the segmentation of the region complete, and moves on by identifying the next seed and growing the next region. The algorithm completes when all the points have been assigned to a region.

In the proposed implementation we use a variant of the above algorithm that also takes into consideration color information, implemented in the `pcl::segmentation::RegionGrowingRGB` class. Practically this differs from the original algorithm in two ways. First, color information is used instead of normals to identify neighboring points. Second, a merging algorithm is ran to control over- and under-segmentation. This algorithm attempts first to merge neighboring clusters depending on their average color difference. Then it uses a user-defined parameter to merge together neighboring clusters that appear to be too small.

The advantage of such algorithm is that it usually performs better on people in terms of precision since a body is not a smooth surface. The merging algorithm is also useful to merge together body parts that are usually identified separately (e.g. attaching a hand to an arm covered by a sleeve).

On the other side, this approach has some drawbacks, the main one being its computational complexity. In theory this problem would be mitigated by the fact that we provide the algorithm with initial seeds, so we should be able to exclude building the majority of segments in a scene. In practice though, even though the implementation found in PCL exposes a method to build a region based on a seed, this method still performs the full scene segmentation first.

One possible solution to this problem could be to just grow a region based on the provided seed and to exclude points that don't match the criteria for the region being grown. The drawback of this is that the step in which adjacent and similar regions are merged is lost, making the algorithm less robust.

5 Implementation details

The effort to realize the algorithm described so far resulted in the implementation of a number of different applications that support the main algorithm. In particular, in the accompanying codebase the source code of the following programs can be found:

- a tool to extract faces from a scene. The tool provides a graphical user interface that allows the user to move around a square viewport centered around the selected point. A point can be selected with **Shift + Click**. The position of the viewport can be adjusted using the arrow keys. Pressing **S** saves the sample to a separate file. The program accepts as input a folder where to read the scenes from, and one where to save the samples to. An example invocation is as follows:

```
make extract_samples in=dataset/samples out=dataset/positive_frontal
```

- an application that trains a cascaded face detector given a set of positive samples and a scene containing only negative samples. The application uses the training procedure described in the preceding paragraphs. The output of the training is a YAML file containing the parameters that can be used to initialize a trained detector. The application can be invoked as follows:

```
make train positive_set=dataset/positive_frontal \
negative_sample=negative.pcd
```

The input consists of two arguments: a folder where the positive samples are located, and the path to a PCD scene that contains only negative samples.

- a sample application that scans a set of scenes for people. The application accepts as input a folder containing PCD scenes. Each scene will be analyzed sequentially.

```
make run in=dataset/samples
```

In addition to this, the codebase contains a set of utilities to facilitate compilation and execution of the applications in a number of different environments. Namely:

- a **CMakeLists.txt** file used by CMake to compile all the applications described above

- a **Makefile** which provides shortcuts for running the applications described above in a virtual environment
- a **Vagrantfile** configuration file and a **provisioning** folder. These are used to set up a Vagrant virtual machine capable of compiling and running the applications. Tasks invoked from the **Makefile** automatically bring up the virtual machine if not already live

The source code of the project is freely available on GitHub.

6 Results

Tests on the implementation of the algorithm described so far provided some results that definitely increase the confidence on the validity of the approach, but overall performance isn't brilliant. The reason for this is to be found in the limited size of the training dataset. Also, despite the speed of the face detector, the speed performance of the segmentation algorithm penalizes the overall time taken by the algorithm to identify a person.

6.1 Training dataset

The training dataset consists in a given set of 40 PCD scenes, one of which being a negative sample not containing any person. Using this scene as a source for negatives yields a negative training set of 2392 samples.

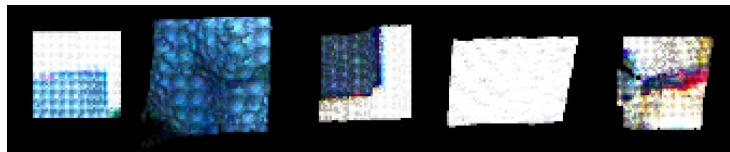


Figure 4: Example of negative samples used in the training process

To build the positive training set, instead, the tool described above has been used. Given the small size of the overall set, heads in all kinds of positions have been selected. Also, the same head has been used multiple times, by slightly shifting the selection window around it. With this approach it was possible to build a positive set of 560 samples.

These numbers are barely enough to obtain solid results in training (especially given that the algorithm is trained to recognize heads in different positions). As a comparison, in their original paper Viola and Jones used a positive training set of around 5000 images, and the pool of negative samples had a size of almost 10000 - out of which a random set of maximum 6000 images was selected for each trained layer.



Figure 5: Example of positive samples used in the training process

In practical terms, this means the algorithm saturates rather quickly and performances can't be improved beyond a certain point without making the training algorithm itself unstable. The final consequence of this is that the final detector has just 3 stages, and in each stage a single detector has been sufficient to reach the best result possible.

6.2 Detection performance

To practically test the validity of the approach, the trained detector has been applied to the set of samples also used for the training itself. It's important to note that this is not an optimal approach, as in these cases it's always recommended to keep the training set and the validation set disjoint to prevent overfitting of the trained classifier. Nevertheless, given the limited amount of samples available, the same set has been used both for training and validation.

The results can be analyzed along two distinct dimensions: precision and accuracy.

We can define precision as the algorithm's capability of identifying a person within a scene, while accuracy deals with how precisely the algorithm identifies a person's boundaries.

When talking about precision, we can use standard measurements like detection rate and false positive rate. To go a bit more in depth, we can also analyze some results and try to characterize the algorithm's performance in more detail.

In terms of detection rate, the results are highly variable as expected from the training step. Performance ranges from 0% to 100%, with an average of 58.57%.

Figures 6 and 7 give a visual clue of this variability. In the first case every person in the scene is correctly detected, while in the second case the only person present is missed by the algorithm. On average detection on a scene looks more like figure 8, where not all the people are correctly identified.



Figure 6: Detection with 100% detection rate

Focusing then on the false positive rate, we can see the results are also variable but a bit more consistent due to the higher amount of negatives present in a scene. Results vary in a range of 0% to 0.06%. Visually, this translates to a few different patterns. A significant example is the output of the detection on the scene containing no positives (fig. 9). While the algorithm rejects most of the samples, it still identifies a person in three of the almost 9600 subwindows.

Another possible case is the one depicted in figure 10 where along with a true positive, a false positive is detected in the chair.

The numbers given so far only considered a detection as either correct or incorrect depending on whether mentioned detection was including part of a person or not. What remains to be analyzed is the accuracy of those detections, meaning how much of a person is detected and how much is instead left out.

The analysis in this case has been performed on a qualitative level, since it's hard to precisely define a person's boundaries so performing an actual measurement brings little value in addition to what a qualitative inspection can give.

Also in this case the results vary in quality. An example of good accuracy is represented by the right detection in figure 11: basically the entire figure is enclosed in the detection cube. At the same time, the detection on the left

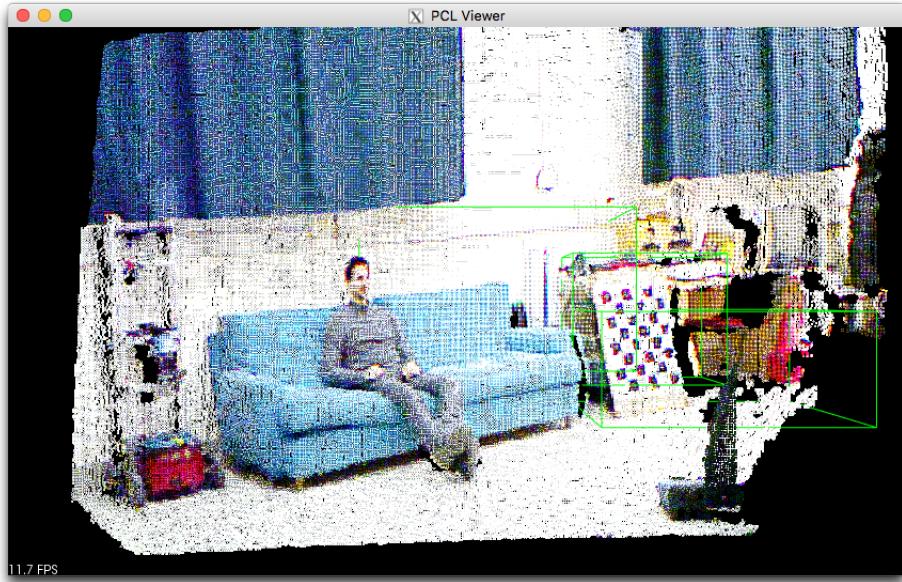


Figure 7: Detection with 0% detection rate

of the scene depicts a common level of low accuracy: the detection cube is centered on the person but doesn't enclose all of it. Another pattern can be seen in figure 12, where the detection cube on the right excluded the lower body part and at the same time appears to be oversized with respect to the actual person.

A last case worth mentioning is that of multiple detections on a single person. The case is easy to understand by simply looking at figure 13: basically a single person is segmented with multiple positive detections. In this case, the detected areas are fundamentally non-overlapping. While overall the whole person is detected, this obviously cannot be considered accurate.

Overall, detection results reflect the results that have been obtained during the training of the detector. They don't make the algorithm "ready for production use", but at the same time they show a clear path for performance improvement: on one side, training the algorithm with a larger training set (and possibly a separate validation set) is likely to give better results in term of precision, while a more precise fine-tuning of the segmentation algorithm could improve the algorithm's accuracy.



Figure 8: The algorithm correctly detects only 3 out of 4 people in the scene

6.3 Speed considerations

All tests have been run in a virtual machine using a MacBook Pro, which hosts a 2.6 GHz Intel Core i5. Due to the significant overhead of running processes inside a virtual machine, the numbers presented here should be considered more in relative terms to each other rather than absolute measures of peak performance.

Tests have been conducted to determine the algorithm's speed. Overall this is the sum of different steps being performed. Some of these steps can be considered as part of the core algorithm function, while others are accessory to its practical implementation and can be excluded. For the sake of completeness all of them have been measured.

The first measured step is opening the scene file. On the test platform this happens on average in around 1800 ms. This step can be made faster if for example the input is provided by a camera, thus having a continuous in-memory stream of scenes (i.e. no disk round-trip to fetch the data).

The second measured step is the face detection on the whole scene. This runs on average in 265 ms. As predicted, this is a very fast operation due to the low amount of computations needed by the cascaded classifier. It's worth noting though that the time complexity is not constant, since the more people in the scene, the more times the classifier will have to run till the end. Also, its speed depends on the amount of stages and features per

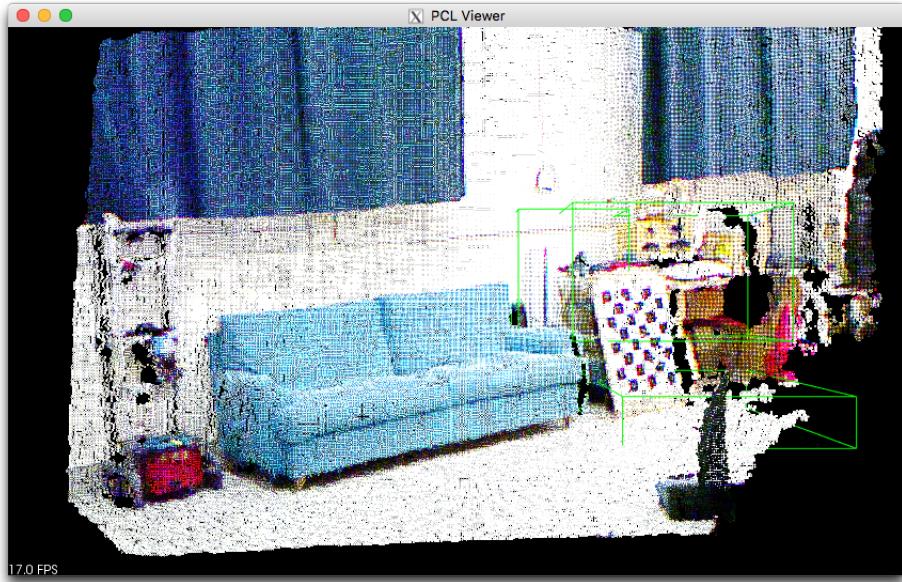


Figure 9: The algorithm correctly detects only 3 out of 4 people in the scene

stage; a classifier trained with a more extended set of samples is going to be more accurate but also slower on average.

The third measured step is the first body segmentation in a scene. This has been measured at 60 s on average. As already discussed, since with the current implementation the whole scene is segmented once, this step results being extremely slow.

The fourth measured step is segmentation of a body after the first one has been identified. This on average runs in 2 ms. As expected, this operation is very fast since the scene has already been segmented, so at this point all that's left to do is to search the cluster containing the centroid of the face that has been identified.

The last measure is made around building the visual output of the algorithm. On average this happens in less than 1 ms. As the first step, this is also an accessory operation, which speed can change depending on the actual implementation and interface requirements; nevertheless in this case its running time is negligible with respect to the rest.

Given the analysis above, one conclusion can be drawn. In its current incarnation, the algorithm is not suitable for real-time detection of people. It is important to note, though, that this is not necessarily due the algorithm's design per-se, but a few other specific components might affect this consequence:

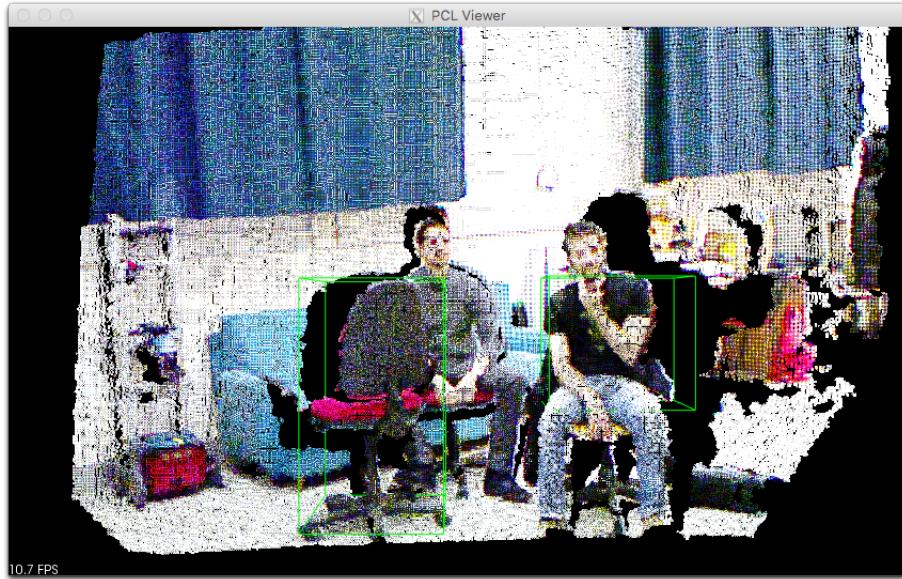


Figure 10: A true detection is missed and a false positive is detected instead

- the actual hardware where the experiment has been run
- the fact that the current implementation doesn't take advantage of a GPU to speed up computation
- the fact that the current implementation is strictly sequential (i.e. it doesn't take advantage of any possible parallelism)

These factors are subject to change in the future and might be subject of a development iteration aimed at improving the algorithm's overall performance.

7 Further developments

The results observed by developing the present paper have been useful not only to validate the proposed approach, but also as inspiration for possible improvements of the base technique.

An obvious improvement, achievable with what built already in this paper, is training the detector with more samples, and possibly with a more diverse set. A fundamental principle of machine learning is that the more data is used for training, the better the results. Also, to make the algorithm suitable for a real-world case it's indispensable for it to be generic enough not to require training every time a scene changes. Using a more diverse

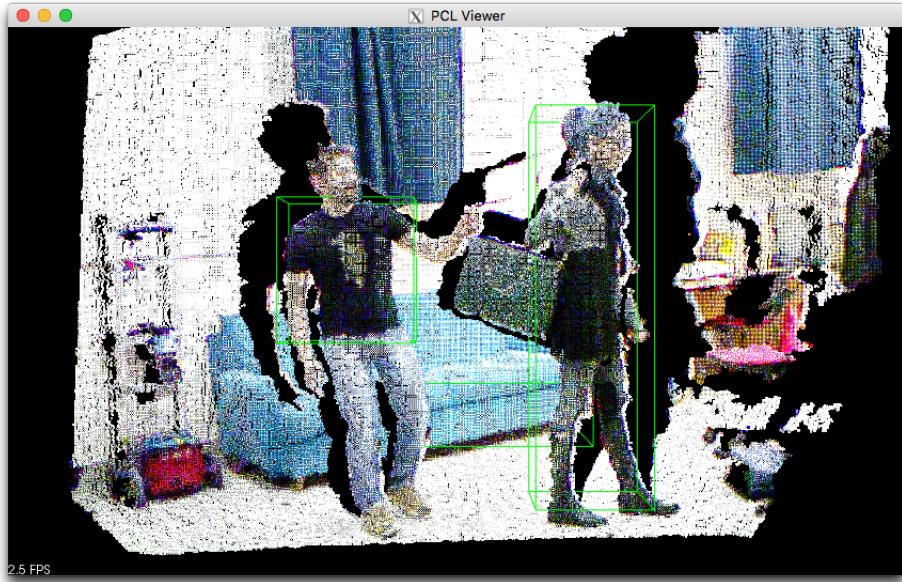


Figure 11: The figure on the right is correctly enclosed in the detection cube, while accuracy for the person on the left is lower

set of samples (e.g. different scenes, different subject) would definitely be of great help.

Talking strictly about improvements over the algorithm itself instead, one first idea is to experiment with different information when calculating features. For example, information about a pixel’s intensity could be substituted with its depth information. The latter might prove to be a more robust information channel rather than the former, since intensity needs normalization to account for varying lighting conditions.

Another possible improvement is based on the observation above, and consists in running multiple detectors in parallel and process the results with a consensus algorithm. The different detectors can be of the same type, trained to detect different poses (e.g. to identify profile faces rather than frontal positions); or they could be of different types, supporting each other in case of difficult detections (e.g. an intensity-based detector running in parallel with a depth-based one, to compensate sensor’s limitations in specific areas).

Another aspect that in this paper hasn’t been considered is that of parallelization. On a high level a cascaded detector isn’t particularly suited for a parallel implementation, but each stage could benefit from that. Since a strong classifier requires computing a certain number of independent features over the same scene, it’s easy to imagine that calculating the feature



Figure 12: The detection on the right is oversized and excludes the lower body part

values in parallel would speed up the algorithm’s overall performance.

A possible improvement might also lie in the use of color information: for example, human skin lies within a specific and restricted range of colors. This fact could be used to speed up removing negatives, so that the final cascaded detector can be trained on more specific cases (incrementing its overall complexity but at the same time making it more refined and precise).

Finally, one aspect that here hasn’t been considered is the use of the algorithm in a sequence of frames. It’s easy to imagine how information from a previously observed scene could help taking better decisions in subsequent frames through subject tracking.

References

- [1] Yoav Freund and Robert E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting, 1996.
- [2] Paul Viola and Michael Jones. Robust real-time face detection. *International Journal of Computer Vision*, 57:137–154, 2004.

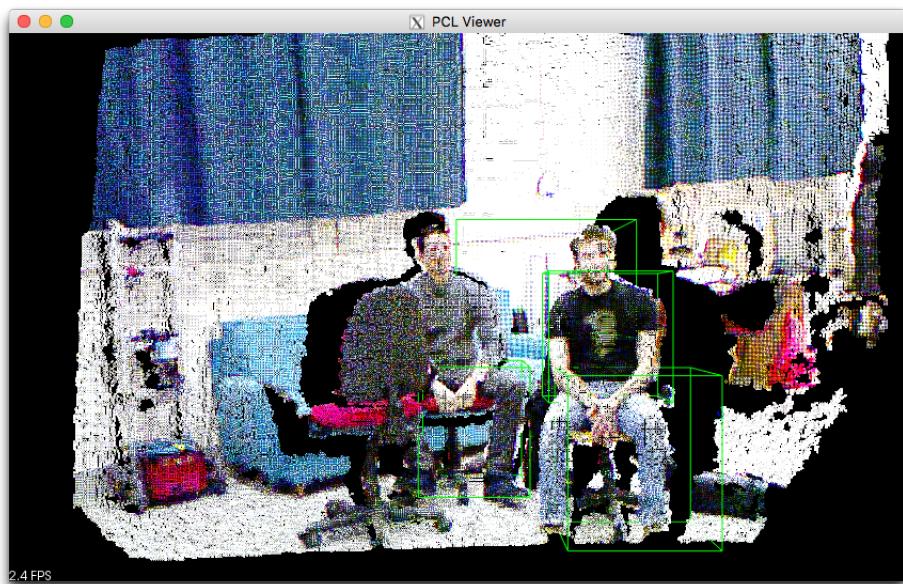


Figure 13: The person on the right is detected twice with non-overlapping areas